

Spam Email Classification

Advanced ML Pipeline with OpenSpec Workflow

Report Generated: 2025■11■14■

Executive Summary

This report documents a complete end-to-end machine learning project for spam email classification. The project implements a logistic regression model trained on 5,574 SMS messages with 96.95% test accuracy. The implementation includes data preprocessing, model training, evaluation, and an interactive Streamlit web application for real-time classification and analysis.

Project Overview

| Aspect | Details |
|------------------|-----------------------------|
| Dataset Size | 5,574 SMS messages |
| Spam Ratio | 13.4% (747 spam, 4,827 ham) |
| Model Type | Logistic Regression |
| Test Accuracy | 96.95% |
| Precision (Spam) | 100% |
| Recall (Spam) | 77.18% |
| F1 Score | 0.871 |
| Vectorization | TF-IDF (max 5,000 features) |
| N-grams | Unigrams and Bigrams (1-2) |

Threshold Sweep Analysis

The following table shows model performance metrics across different decision thresholds, allowing optimization for specific use cases (prioritize precision or recall):

| Threshold | Precision | Recall | F1 Score |
|-----------|-----------|--------|----------|
| 0.1 | 0.6558 | 0.9920 | 0.7896 |
| 0.2 | 0.9283 | 0.9705 | 0.9490 |
| 0.3 | 0.9744 | 0.9170 | 0.9448 |
| 0.4 | 0.9893 | 0.8648 | 0.9229 |

| | | | |
|-----|--------|--------|--------|
| 0.5 | 0.9915 | 0.7831 | 0.8751 |
| 0.6 | 0.9910 | 0.5877 | 0.7378 |
| 0.7 | 1.0000 | 0.3548 | 0.5237 |
| 0.8 | 1.0000 | 0.1299 | 0.2299 |
| 0.9 | 1.0000 | 0.0147 | 0.0290 |

Data Preprocessing Pipeline

The project implements a comprehensive 7-stage text preprocessing pipeline:

| Stage | Operation | Purpose |
|-----------------------------|---------------------------|--------------------|
| 1. Raw | Original text | Baseline reference |
| 2. Lowercase | Convert to lowercase | Normalization |
| 3. Contact Masking | Mask emails/phones | Remove PII |
| 4. Number Replacement | Replace digits with <NUM> | Generalization |
| 5. Punctuation Removal | Remove special characters | Simplification |
| 6. Whitespace Normalization | Normalize spaces | Formatting |
| 7. Stopword Removal | Remove common words | Feature reduction |

Key Features

- 1. Multi-Format CSV Support:** The application supports both simple 2-column CSV format and advanced 9-column preprocessing pipeline format for detailed text transformation analysis.
- 2. Interactive Dashboard:** Streamlit-based web application with real-time classification, token analysis, and model performance visualization.
- 3. Advanced Analytics:** Threshold sweep analysis, ROC curves, confusion matrices, and precision-recall curves for comprehensive model evaluation.
- 4. CLI Tools:** Command-line utilities for batch prediction, visualization generation, and model training with custom parameters.
- 5. Professional Documentation:** Comprehensive README, quick-start guides, and delivery summaries with usage examples and technical details.

Technology Stack

| Component | Technology | Purpose |
|-----------------|-----------------------------|----------------------------------|
| Language | Python 3.12+ | Core implementation |
| ML Framework | Scikit-learn | Model training & evaluation |
| Data Processing | Pandas, NumPy | Data manipulation |
| Visualization | Plotly, Matplotlib, Seaborn | Interactive & publication charts |
| Web Framework | Streamlit | Interactive dashboard |
| Serialization | joblib | Model & vectorizer storage |
| Deployment | Streamlit Cloud | Public web application |
| Version Control | Git, GitHub | Code management |
| Workflow | OpenSpec | Specification-driven development |

Project Structure

```
■■■ app.py # Streamlit web application
■■■ train.py # Model training script
■■■ requirements.txt # Python dependencies
■■■ README.md # Project documentation
■■■ src/
■ ■■■ data_loader.py # Data loading & preprocessing
■ ■■■ model_trainer.py # Model training & evaluation
■■■ scripts/
■ ■■■ predict_spam.py # CLI prediction tool
■ ■■■ visualize_spam.py # Visualization toolkit
■ ■■■ generate_report.py # PDF report generation
■■■ data/
■ ■■■ sms_spam_clean.csv # Clean 2-column format
■ ■■■ sms_spam_preprocessing.csv# 9-column preprocessing pipeline
■ ■■■ sms_spam_no_header.csv # Original format
■■■ models/
■ ■■■ logistic_regression.pkl # Trained model (joblib)
■ ■■■ vectorizer.pkl # TF-IDF vectorizer
■ ■■■ label_mapping.json # Label mappings
■ ■■■ metrics_logistic_regression.json # Performance metrics
■ ■■■ threshold_sweep.json # Threshold analysis
■ ■■■ test_predictions.json # Test predictions for ROC
■■■ docs/
■■■ PREPROCESSING.md # Preprocessing documentation
```

Model Performance Results

The Logistic Regression model achieved excellent performance on the spam classification task:

| Metric | Value | Description |
|--------------------|--------|--|
| Test Accuracy | 96.95% | Overall correctness of predictions |
| Precision (Spam) | 100% | All spam predictions were correct |
| Recall (Spam) | 77.18% | 77% of actual spam was detected |
| F1 Score | 0.871 | Harmonic mean of precision & recall |
| ROC-AUC | ~0.98 | Excellent discriminative ability |
| Specificity | 100% | No false positive rate |
| True Negative Rate | 100% | All legitimate emails correctly classified |

How to Use

1. Running the Web Application:

```
streamlit run app.py
```

Open your browser to <http://localhost:8501>

2. Making Predictions (CLI):

```
python scripts/predict_spam.py --text "Your message here"
```

3. Batch Predictions:

```
python scripts/predict_spam.py --input data.csv --output predictions.csv
```

4. Generating Visualizations:

```
python scripts/visualize_spam.py --input data.csv --dist --tokens
```

5. Training Model:

```
python train.py --model logistic_regression
```

Conclusions & Future Work

Achievements:

- Successfully built a high-accuracy spam classification model (96.95% accuracy)
- Implemented comprehensive data preprocessing pipeline with 7 stages
- Created professional interactive dashboard with real-time classification
- Developed CLI tools for batch processing and automation
- Demonstrated OpenSpec specification-driven development workflow

Future Enhancements:

- Support for multiple languages and character sets
- Ensemble models combining multiple algorithms
- Active learning with user feedback integration
- Advanced NLP techniques (BERT, transformers)
- Deployment to cloud platforms (AWS, GCP, Azure)
- Real-time model retraining with new data

